

CSE 544: Homework 5

Some Theory Question

Dan Suci

Released: February 27, 2021
Hard Deadline: Monday, March 15, 2021

Name: _____

Question	Points	Score
1	40	
2	40	
3	20	
Total:	100	

Write your answers in a file (pdf, or docx, or txt, or md) in the **submission** directory.
Write the problem number and answer, for example (not real answers):

2.b.(ii) 0.73

3.a: ABORT ABORT ABORT ABORT

Don't spend time trying to make the file look very nice, but make sure your answers are clear and unambiguous.

1. (40points)

Consider a database schema with three relations and their sizes:

```
Likes(drinker, beer)      // 100,000 tuples
Frequents(drinker, bar)   // 2,000,000 tuples
Serves(bar, beer)         // 5,000,000 tuples
```

Consider the following two queries:

Q1:

```
select *
from Likes, Frequents, Serves
where Likes.drinker = Frequents.drinker
      and Frequents.bar = Serves.bar
```

Q2:

```
select *
from Likes, Frequents, Serves
where Likes.drinker = Frequents.drinker
      and Frequents.bar = Serves.bar
      and Likes.beer = Serves.beer
```

Answer the questions below:

- (a) (20 points) What is the largest possible output of Q1?
- (b) (20 points) What is the largest possible output of Q2?

2. (40points)

Consider the relation **Products**(*pid*, *name*, *category*, *price*), with $N = 1,000,000,000$ tuples and $V(\text{Product}, \text{category}) = 100,000$ distinct categories.

For the questions below use a calculator and write an approximate answer, or choose from the available options (whichever you prefer):

- (a) We hash-partition **Products** by the key *pid* over $P = 1,000$ servers.
- (10 points) What is the probability that the load of some server exceeds the expected load by more than 50%?
 - ≥ 0.99 (almost certain)
 - $(0.5, 0.99)$
 - $(0.1, 0.5)$
 - $(0.01, 0.1)$
 - $(0.001, 0.01)$
 - < 0.001 (absolutely tiny)
 - (10 points) What is the probability that the load of some server exceeds the expected load by more than 0.5%?
 - ≥ 0.99 (almost certain)
 - $(0.5, 0.99)$
 - $(0.1, 0.5)$
 - $(0.01, 0.1)$
 - $(0.001, 0.01)$
 - < 0.001 (absolutely tiny)
- (b) Now suppose we hash-partition **Products** by *category* over $P = 1,000$ servers. Assume that each category occurs the same number of times.
- (10 points) What is the probability that the load of some server exceeds the expected load by more than 50%?
 - ≥ 0.99 (almost certain)
 - $(0.5, 0.99)$
 - $(0.1, 0.5)$
 - $(0.01, 0.1)$
 - $(0.001, 0.01)$
 - < 0.001 (absolutely tiny)
 - (10 points) What is the probability that the load of some server exceeds the expected load by more than 0.5%?
 - ≥ 0.99 (almost certain)
 - $(0.5, 0.99)$
 - $(0.1, 0.5)$
 - $(0.01, 0.1)$
 - $(0.001, 0.01)$
 - < 0.001 (absolutely tiny)

3. (20points)

- (a) (20 points) Consider a multiversion concurrency control manager. Indicate what happens for each question mark below. If it is a write operation say whether it is executed or aborted; if it is a read operation, indicate the value returned (e.g. say “ $X = 9$ ”), or say that it is aborted. An instruction like $W_t(X = v)$ means that the transaction with timestamp t writes the element X , and the value of X being written is v .

START ₁	(executed OK)
$W_1(X = 1)$	(executed OK)
COMMIT ₁	(executed OK)
START ₂	(executed OK)
START ₃	(executed OK)
START ₄	(executed OK)
$W_4(X = 4)$	(executed OK)
COMMIT ₄	(executed OK)
START ₅	(executed OK)
START ₆	(executed OK)
START ₇	(executed OK)
$W_7(X = 7)$	(executed OK)
COMMIT ₇	(executed OK)
$R_5(X)$?
$R_3(X)$?
$W_2(X = 2)$?
$W_6(X = 6)$?